



Disparities in Air Pollution Exposure

ZHUOPEI JIN | DANIEL SUNG | ALYCE KEYS



- 01 Background**
- 02 Problem Statement**
- 03 Method and Data**
- 04 Results and Discussion**
- 05 Conclusion**



Background

Air quality is the measure of how polluted the air is in a particular area

Air pollution is a major environmental risk to public health and can result in health issues including heart disease, lung and bronchus cancer, stroke, and death


PM2.5 are fine particles less than 2.5 micrometers in diameter that pose risk to health and visibility

Problem Statement

It is well documented that racial/ethnic minorities and people of low socioeconomic status in the US are at a higher risk of death from being exposed to PM2.5 (Di, Q. et al. 2017)

We aim to explore how different demographics are affected by PM2.5 in their respective regions in Texas

We then plan to use these models to predict



Method and Data

ZCTA

**Demographic
Data**

**Environmental
Data**



Method and Data

ZCTA

ZCTA: ZIP Code Tabulation Area

Effectively captured population patterns using tracts

Unnamed: 0 year zcta pm25

| | | | | |
|--------|--------|------|---------|-----|
| 789255 | 789256 | 2014 | 99921.0 | NaN |
| 789256 | 789257 | 2015 | 99921.0 | NaN |
| 789257 | 789258 | 2016 | 99921.0 | NaN |
| 789258 | 789259 | 2017 | NaN | NaN |
| 789259 | 789260 | 2018 | NaN | NaN |

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 789260 entries, 0 to 789259
Data columns (total 3 columns):
 # Column Non-Null Count Dtype

 0 year 789260 non-null int64
 1 zcta 686681 non-null Int64
 2 pm25 689745 non-null float64
 dtypes: Int64(1), float64(1), int64(1)
 memory usage: 18.8 MB

Method and Data

Demographic Data

Social Explorer, ACS(5 years) data from 2011 to 2021

- Total population
- Race
- Median Household Income
- Education Attainment

#ACS Yearly Data

```
acs_2011 = pd.read_csv('ACS 2011.csv')
acs_2011['year'] = 2011
acs_2012 = pd.read_csv('ACS 2012.csv')
acs_2012['year'] = 2012
acs_2013 = pd.read_csv('ACS 2013.csv')
acs_2013['year'] = 2013
acs_2014 = pd.read_csv('ACS 2014.csv')
acs_2014['year'] = 2014
acs_2015 = pd.read_csv('ACS 2015.csv')
acs_2015['year'] = 2015
acs_2016 = pd.read_csv('ACS 2016.csv')
acs_2016['year'] = 2016
acs_2017 = pd.read_csv('ACS 2017.csv')
acs_2017['year'] = 2017
acs_2018 = pd.read_csv('ACS 2018.csv')
acs_2018['year'] = 2018
acs_2019 = pd.read_csv('ACS 2019.csv')
acs_2019['year'] = 2019
acs_2020 = pd.read_csv('ACS 2020.csv')
acs_2020['year'] = 2020
acs_2021 = pd.read_csv('ACS 2021.csv')
acs_2021['year'] = 2021
```

```
merged_acs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 346566 entries, 0 to 364973
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Geo_FIPS        346566 non-null   int64  
 1   Geo_GEOID       346566 non-null   object  
 2   zcta            346566 non-null   int32  
 3   Median Household Income 346566 non-null   float64 
 4   total population 346566 non-null   int64  
 5   population Density 346566 non-null   float64 
 6   white_pop       346566 non-null   int64  
 7   black_pop       346566 non-null   int64  
 8   asian_pop       346566 non-null   int64  
 9   hispanic_pop    346566 non-null   int64  
 10  year             346566 non-null   int64  
 11  Bachelor's degree or higher 346566 non-null   int64  
dtypes: float64(2), int32(1), int64(8), object(1)
memory usage: 33.1+ MB
```

Method and Data


Environmental Data

- PM2.5 Concentrations (Target Variable)
- Traffic Volume


traffic

| | zcta19 | year | intp_flag | i_mean_traffic | i_min_traffic | i_max_traffic | i_count_traffic | i_total_traffic | i_mean_hw_traffic | i_min_hw_traffic | ... | i_total_hw_traffic |
|--------|--------|------|-----------|----------------|---------------|---------------|-----------------|-----------------|-------------------|------------------|-----|--------------------|
| 0 | 1001 | 1993 | 0 | 25872.0000 | 25872.0000 | 25872.000 | 1.0 | 25872.00 | NaN | NaN | ... | 0. |
| 1 | 1001 | 1994 | 0 | 26000.0000 | 26000.0000 | 26000.000 | 1.0 | 26000.00 | NaN | NaN | ... | 0. |
| 2 | 1001 | 1995 | 0 | 15700.0000 | 4500.0000 | 23000.000 | 5.0 | 78500.00 | NaN | NaN | ... | 0. |
| 3 | 1001 | 1996 | 0 | 40600.0000 | 40600.0000 | 40600.000 | 1.0 | 40600.00 | NaN | NaN | ... | 0. |
| 4 | 1001 | 1997 | 0 | 32466.6700 | 9400.0000 | 56600.000 | 3.0 | 97400.00 | NaN | NaN | ... | 0. |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 675321 | 99929 | 2014 | 1 | 891.6190 | 331.3333 | 1852.667 | 9.0 | 7648.00 | NaN | NaN | ... | 0. |
| 675322 | 99929 | 2015 | 1 | 899.4643 | 371.5000 | 1801.750 | 7.5 | 6428.25 | NaN | NaN | ... | 0. |
| 675323 | 99929 | 2016 | 1 | 907.3095 | 411.6667 | 1750.833 | 6.0 | 5208.50 | NaN | NaN | ... | 0. |
| 675324 | 99929 | 2017 | 1 | 915.1548 | 451.8333 | 1699.917 | 4.5 | 3988.75 | NaN | NaN | ... | 0. |
| 675325 | 99929 | 2018 | 0 | 923.0000 | 492.0000 | 1649.000 | 3.0 | 2769.00 | NaN | NaN | ... | 0. |

675326 rows × 22 columns



Research Methodology



Regression Model

Linear regression

Lasso Regression

Random Forest
Regression &
XGBoost

Gradient Boosting
Regressor

| OLS Regression Results | | | | | | |
|-----------------------------|------------------|---------------------|-----------|-------|-----------|-----------|
| Dep. Variable: | pm25 | R-squared: | 0.413 | | | |
| Model: | OLS | Adj. R-squared: | 0.412 | | | |
| Method: | Least Squares | F-statistic: | 352.5 | | | |
| Date: | Tue, 05 Dec 2023 | Prob (F-statistic): | 0.00 | | | |
| Time: | 12:34:04 | Log-Likelihood: | -7675.5 | | | |
| No. Observations: | 4516 | AIC: | 1.537e+04 | | | |
| Df Residuals: | 4506 | BIC: | 1.544e+04 | | | |
| Df Model: | 9 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| | | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| const | 6.7805 | 0.072 | 94.347 | 0.000 | 6.640 | 6.921 |
| i_total_traffic | 4.96e-07 | 5.76e-08 | 8.612 | 0.000 | 3.83e-07 | 6.09e-07 |
| Median Household Income | 2.234e-06 | 1.4e-06 | 1.601 | 0.110 | -5.02e-07 | 4.97e-06 |
| total population | -0.0005 | 9.37e-05 | -5.378 | 0.000 | -0.001 | -0.000 |
| population Density | 0.0004 | 1.58e-05 | 28.397 | 0.000 | 0.000 | 0.000 |
| white_pop | 0.0005 | 9.66e-05 | 5.659 | 0.000 | 0.000 | 0.001 |
| black_pop | 0.0006 | 9.54e-05 | 5.838 | 0.000 | 0.000 | 0.001 |
| asian_pop | 0.0005 | 0.000 | 5.123 | 0.000 | 0.000 | 0.001 |
| hispanic_pop | 0.0005 | 9.4e-05 | 5.631 | 0.000 | 0.000 | 0.001 |
| Bachelor's degree or higher | -3.437e-05 | 8.54e-06 | -4.027 | 0.000 | -5.11e-05 | -1.76e-05 |
| | | | | | | |
| Omnibus: | 128.353 | Durbin-Watson: | 2.016 | | | |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 247.123 | | | |
| Skew: | 0.203 | Prob(JB): | 2.18e-54 | | | |
| Kurtosis: | 4.071 | Cond. No. | 1.86e+06 | | | |
| | | | | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.86e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Model Overview:

R-squared of 41.0% indicates the model explains variability in pm25.

The overall model is statistically significant.

Predictor Variable Analysis

Every Daily Traffic Counts Positive

A one-unit increase is associated with a 5.151e-07 increase in pm25

Median Household Income Positive

Coefficient is positive but not statistically significant ($p = 0.076$).

Total Population Negative

A one-unit increase is associated with a 5.151e-07 increase in pm25

Population Density Positive

A one-unit increase is associated with a 0.0004 increase in pm25

White, Black, Asian, Hispanic Population Positive

Each one-unit increase is associated with an increase in pm25 by the respective coefficients.

Bachelor's Degree or Higher Negative

A one-unit increase is associated with a decrease in pm25 by -2.919e-05.

Regression Model

Linear regression

Lasso Regression

Random Forest
Regression &
XGBoost

Gradient Boosting
Regressor

| | Feature | Coefficient |
|---|-----------------------------|---------------|
| 0 | const | 0.000000e+00 |
| 1 | i_total_traffic | 5.266859e-07 |
| 2 | Median Household Income | 2.324231e-06 |
| 3 | total population | -2.438346e-05 |
| 4 | population Density | 4.387320e-04 |
| 5 | white_pop | 5.281370e-05 |
| 6 | black_pop | 6.950953e-05 |
| 7 | asian_pop | 1.798520e-05 |
| 8 | hispanic_pop | 4.831725e-05 |
| 9 | Bachelor's degree or higher | -3.155886e-05 |

Predictor Variable Analysis

Every Daily Traffic Counts Positive

Median Household Income Positive

Total Population Negative

Population Density Positive

White
Black
Asian
Hispanic Positive

Bachelor's Degree or
Higher Negative

Regression Model

Linear regression

Lasso Regression

Random Forest
Regression &
XGBoost

Gradient Boosting
Regressor

Random Forest Regression & XGBoost

```
: from sklearn.ensemble import RandomForestRegressor  
  
#create regressor object  
model = RandomForestRegressor(n_estimators= 100, random_state= 0)  
  
#fit the regressor with x and y data  
model.fit(x_train, y_train)  
  
: RandomForestRegressor(random_state=0)  
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.  
  
: y_pred= model.predict(X_test)  
  
: print(model.score(X_train, y_train))  
print(model.score(X_test, y_test))  
  
0.953771167771093  
0.7887603072085864
```

```
import xgboost as xg  
from sklearn.metrics import mean_squared_error as MSE  
  
#Instantiation  
  
xgb_r = xg.XGBRegressor(objective= 'reg:linear', n_estimators= 100, seed = 123)  
  
xgb_r.fit(X_train, y_train)  
  
y_pred_xgb = xgb_r.predict(X_test)  
  
print(xgb_r.score(X_train, y_train))  
print(xgb_r.score(X_test, y_test))  
  
0.9495909439186249  
0.756339597268104
```

Regression Model

Linear regression

Lasso Regression

Random Forest
Regression &
XGBoost

Gradient Boosting
Regressor

Gradient Boosting Regressor

```
from sklearn.ensemble import GradientBoostingRegressor

# Instantiate Gradient Boosting Regressor
gbr = GradientBoostingRegressor(n_estimators = 100, max_depth = 1)

# Fit to training set
gbr.fit(X_train, y_train)

# Predict on test set
pred_y = gbr.predict(X_test)


# test set RMSE
test_rmse = MSE(y_test, pred_y) ** (1 / 2)

print(gbr.score(X_train, y_train))
print(gbr.score(X_test, y_test))

0.5146372829636321
0.4992482908383168
```

```
raise, message: 'Failed to read email file', error)
data = data.replace("##Title##", "Confirm Account")
data = data.replace("##Message##", "Hello, User!
```

Principal Component Analysis



Pipeline

With PCA

```
: for i,model in enumerate(pipelines):
    print("{} Test Accuracy: {}".format(pipe_dict[i], model.score(X_test,y_test)))
```

```
Linear Regression Test Accuracy: 0.3524967527659876
Decision Tree Regressor Test Accuracy: 0.2773010752709888
Random Forest Regressor Test Accuracy: 0.5729584606990209
Support Vector Regressor Test Accuracy: 0.39831525573017257
Gradient Boosting Regressor Test Accuracy: 0.4837166608453629
```

Without PCA

```
# Pipeline for: decision tree, random forest regressor, SVC, GradientBoosting Regressor
for pipe in pipelines:
    pipe.fit(X_train, y_train)
for i,model in enumerate(pipelines):
    print("{} Test Accuracy: {}".format(pipe_dict[i], model.score(X_test,y_test)))
```

```
Linear Regression Test Accuracy: 0.39919124655557614
Decision Tree Regressor Test Accuracy: 0.6195287179347277
Random Forest Regressor Test Accuracy: 0.7828656586754518
Support Vector Regressor Test Accuracy: 0.5185248684004499
Gradient Boosting Regressor Test Accuracy: 0.6397047360077701
```

GridSearchCV

```
In [83]: # Specify hyperparameters for the RandomForestRegressor. Testing with different values for hyperparameters
parameters = {
    'regressor__max_features': ['sqrt', 'log2', None],
    'regressor__max_leaf_nodes': [10, 100, 200, None],
    'regressor__n_estimators': [10, 100, 500, 1000],
    'regressor__bootstrap': [True, False],
    'scaler': [StandardScaler(), MinMaxScaler(), Normalizer(), MaxAbsScaler()]
}
```

```
In [84]: grid = GridSearchCV(pipeline_rf, parameters, cv=2).fit(X_train, y_train)
```

```
In [85]: print('Training set score: ' + str(grid.score(X_train, y_train)))
print('Test set score: ' + str(grid.score(X_test, y_test)))
```

```
Training set score: 0.9547499011962798
Test set score: 0.7931548759013509
```

```
In [86]: #Access the best set of parameters
best_params= grid.best_params_
print(best_params)

best_pipe= grid.best_estimator_
print(best_pipe)

{'regressor__bootstrap': True, 'regressor__max_features': 'sqrt', 'regressor__max_leaf_nodes': None, 'regressor__n_estimators': 500, 'scaler': StandardScaler()}
Pipeline(steps=[('scaler', StandardScaler()), ('selector', VarianceThreshold()),
               ('regressor',
                RandomForestRegressor(max_features='sqrt', n_estimators=500))])
```

Conclusion

In Summary, we used ZCTA, ACS, and traffic data to analyze, compare, and correlate with the goal of establishing a prediction system on how pollution (PM2.5) relates to traffic and demographic information.

We found that Random Forest Regression hyperparameters provided the best values

Improvements

To further improve this experiment, we could add more parameters to the gridsearch

We acknowledge several limitations presented by this experiment

Questions?